



Hortonworks Data Platform

November 2015

Guilherme Bracciali
Solutions Engineer - LATAM
gbracciali@hortonworks.com
Skype: gbracciali
+55 11 9 8134 1618

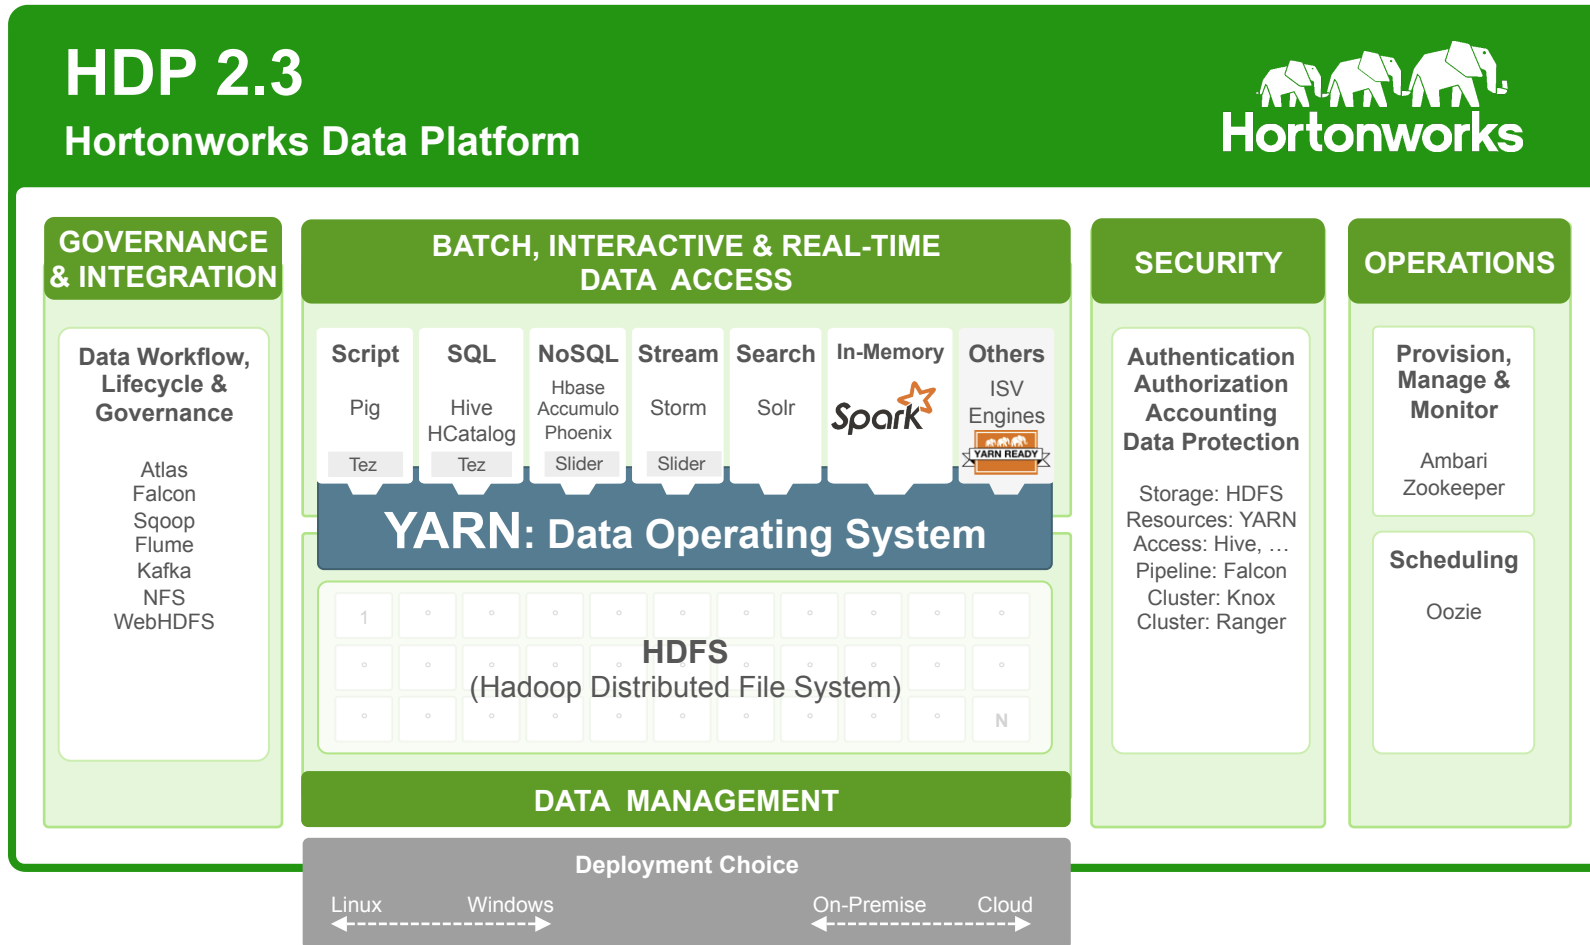
Agenda – Meetup Machine Learning – 05/11/2015

- ✓ Introdução – 20 minutos
 - ✓ Hortonworks
 - ✓ Spark
 - ✓ Zeppelin
- ✓ Demo – Parte 1 – 10 minutos
- ✓ Spark Streaming – 10 minutos
- ✓ Demo – Parte 2 – 10 minutos
- ✓ Apoio – 30 minutos
- ✓ Dúvidas – 10 minutos

Hortonworks

HDP delivers a completely open data platform

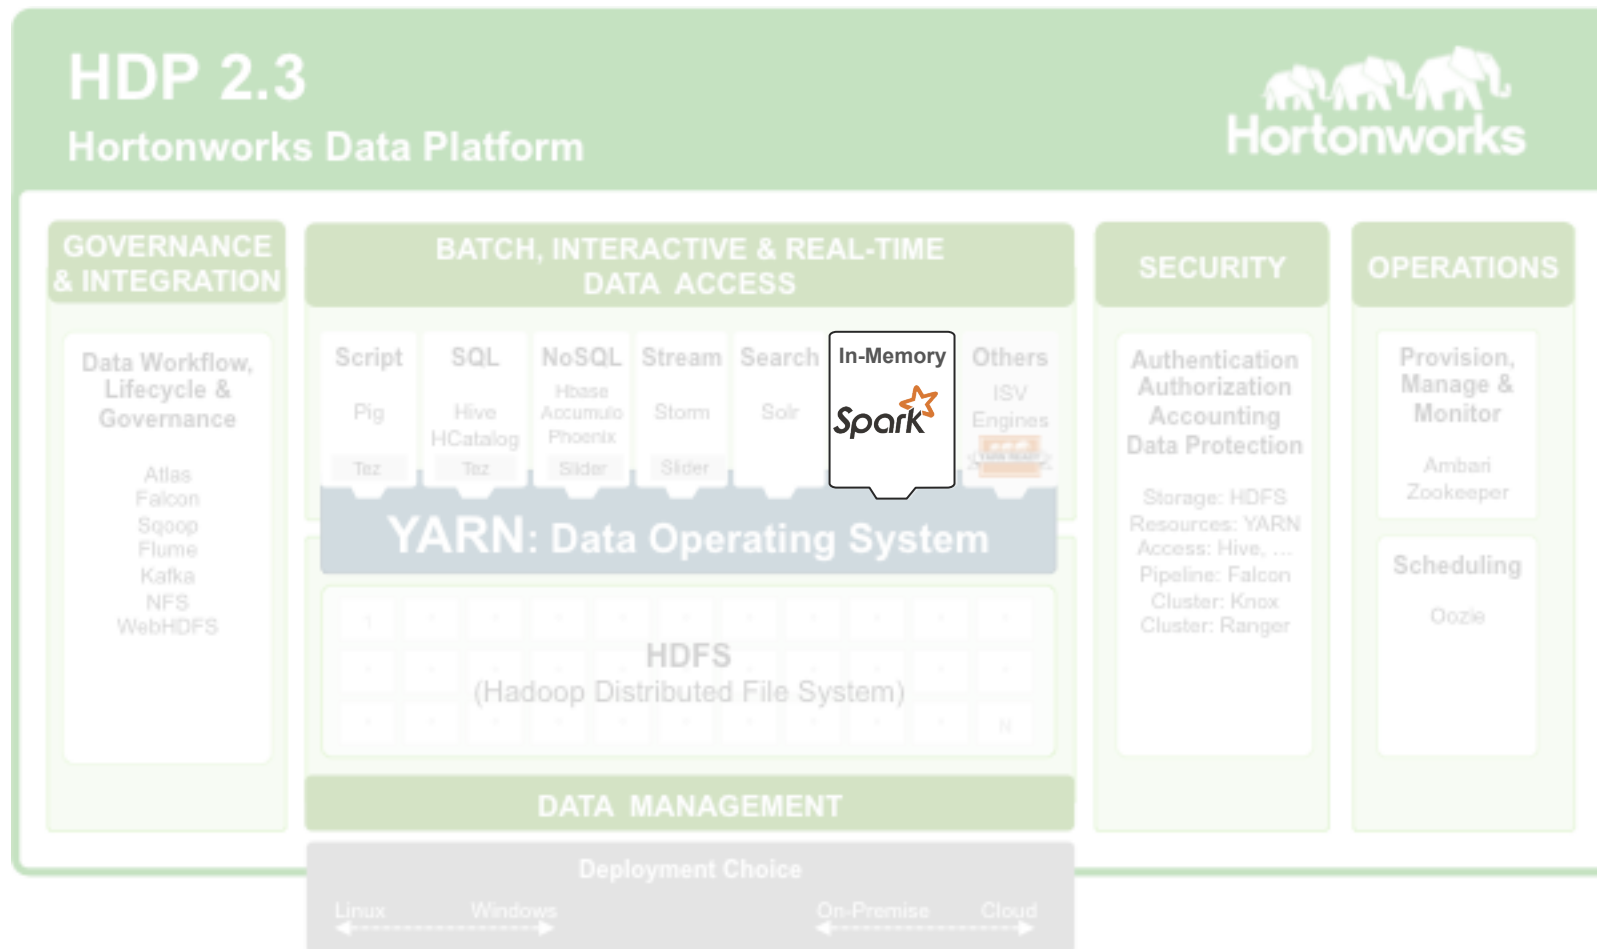
Hortonworks Data Platform provides Hadoop for the Enterprise: a centralized architecture of core enterprise services, for any application and any data.



Completely Open

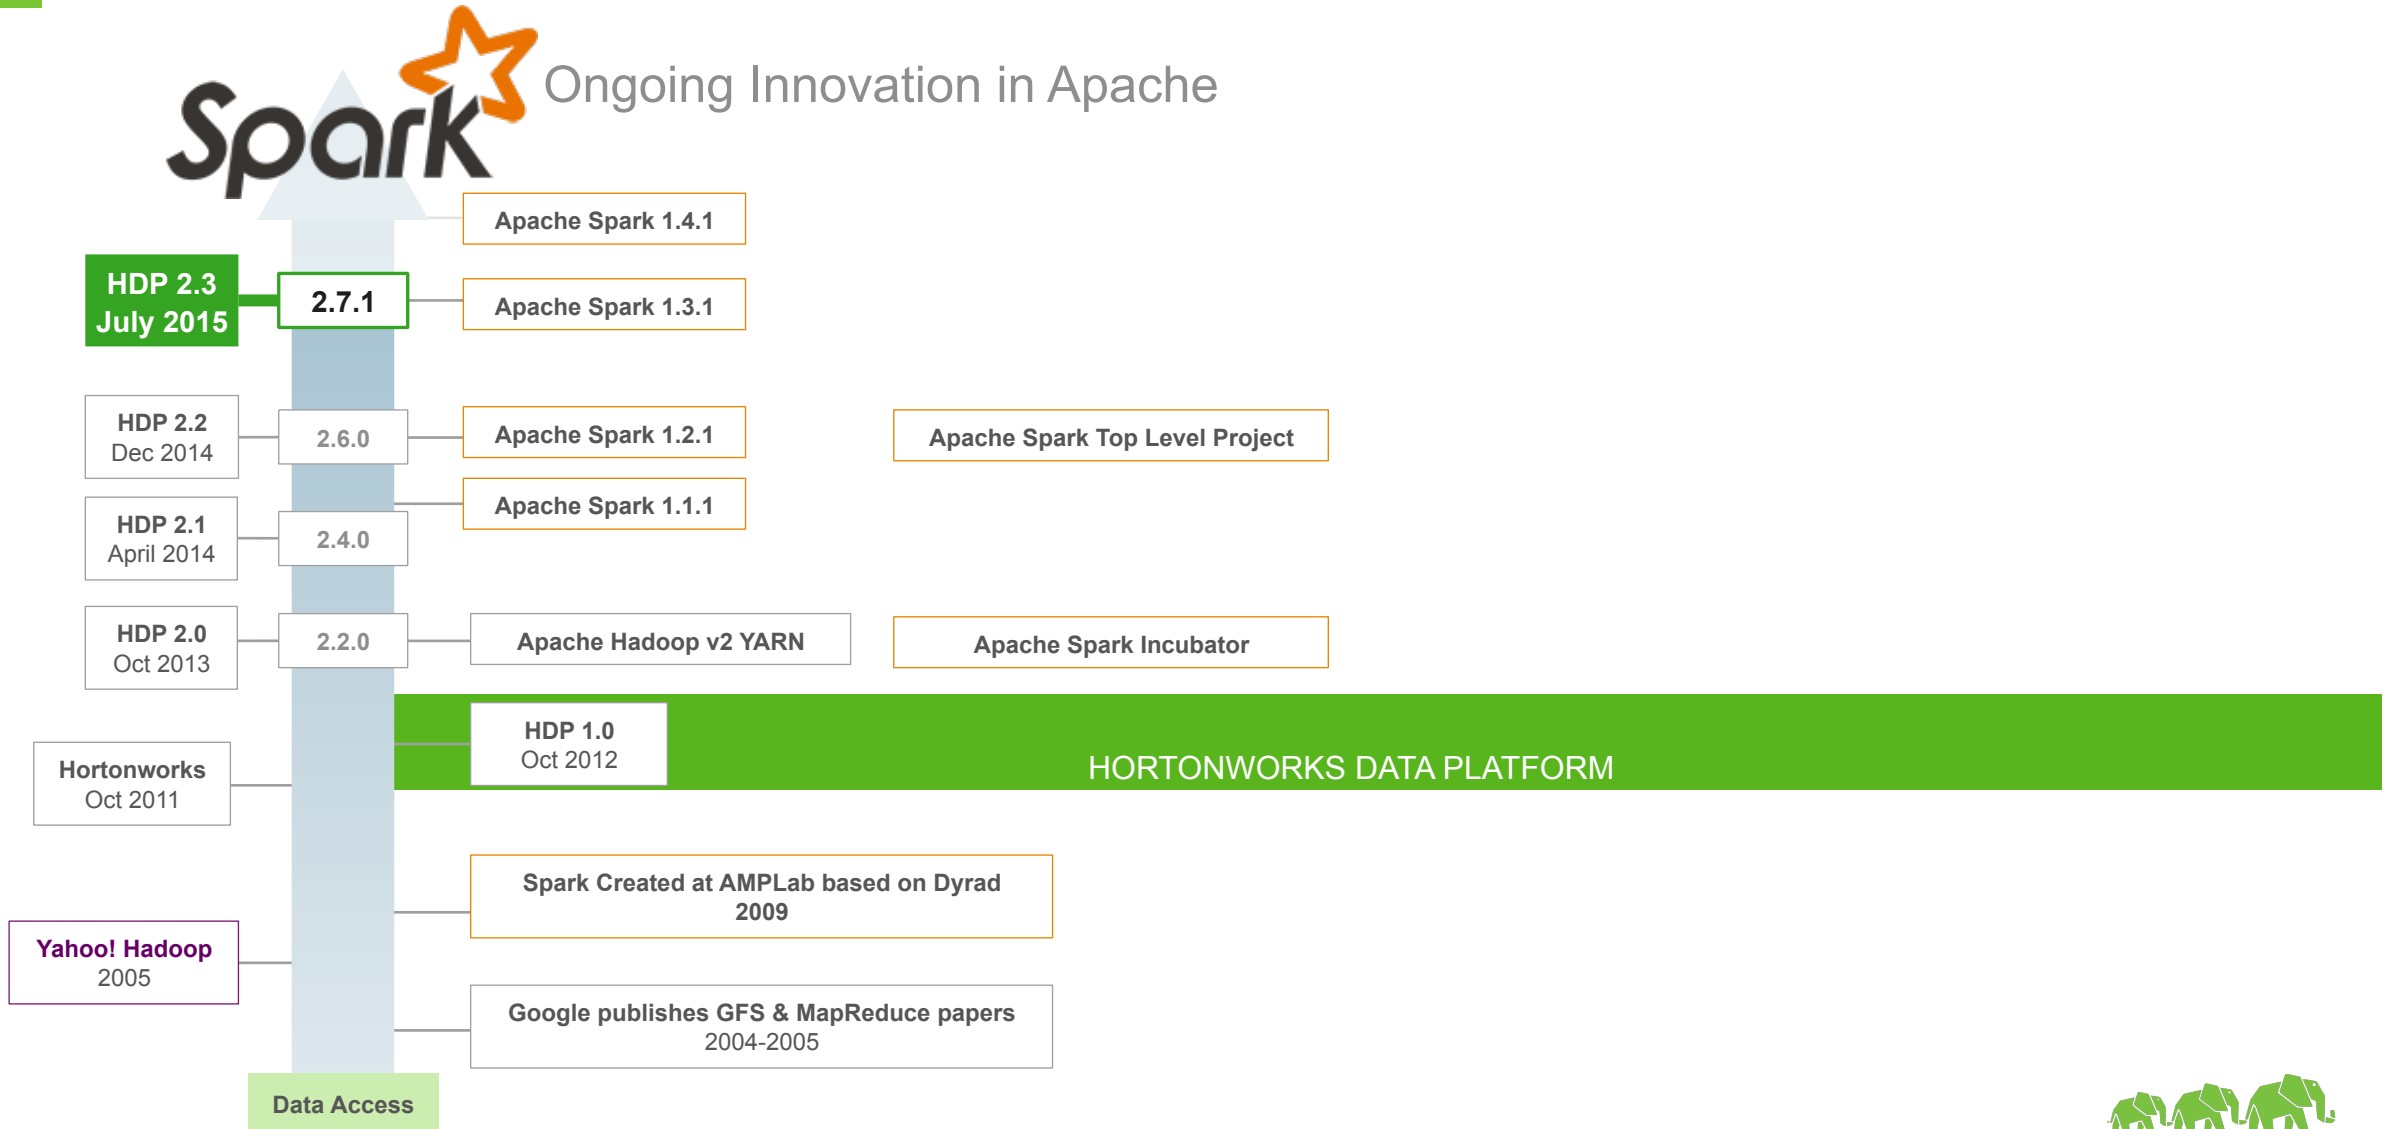
- HDP incorporates every element required of an enterprise data platform: data storage, data access, governance, security, operations
- All components are developed in open source and then rigorously tested, certified, and delivered as an integrated open source platform that's easy to consume and use by the enterprise and ecosystem.

Spark



Spark

Evolution of Apache Spark



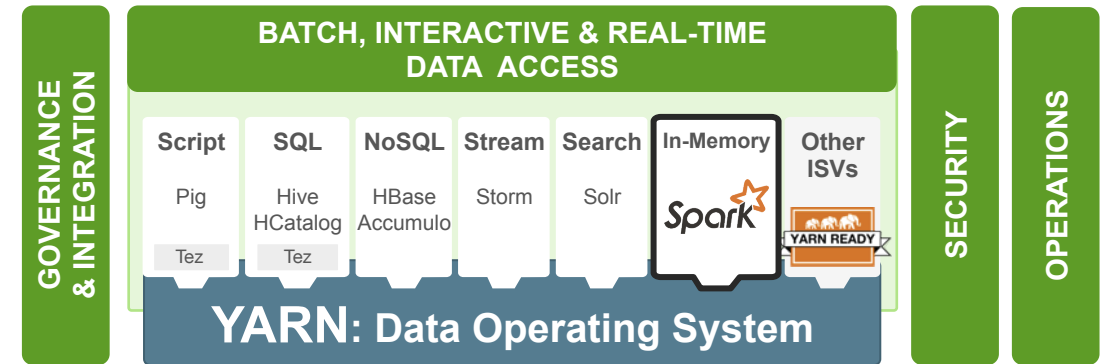
What is Apache Spark?

- Spark is top level Apache Project since February 2014, originally a graduate project at UC Berkeley's AMPLab
- Spark is a general-purpose engine for ad-hoc interactive analytics, iterative machine-learning, and other use cases well-suited to interactive, in-memory data processing of GB to TB sized datasets.
- Spark loads data into memory so it can be queried repeatedly. It can create a “shadow” of data that can be used in the next iteration of a query
- Spark provides simple APIs for data scientists and engineers familiar with Scala (programming language) to build applications
- Spark is built on HDFS
- **Spark is YARN enabled**



Hortonworks Commitment to Spark

Hortonworks is focused on making Apache Spark enterprise ready so you can depend on it for mission critical applications



1. YARN enable Spark to co-exist with other engines

We have already declared it “YARN Ready” so its memory & CPU intensive apps can work with predictable performance along side other engines all on the same set(s) of data.

2. Extend Spark with enterprise capabilities

Ensure Spark can be managed, secured and governed all via a single set of frameworks to ensure consistency. Ensure reliability and quality of service of Spark along side other engines.

3. Actively contribute within the open community

As with everything we do at Hortonworks we work entirely within the open community across Spark and all related projects to improve this key Hadoop technology.

Why We Love Spark at Hortonworks

Made for Data Science

All apps need to get predictive at scale and fine granularity

Democratizes Machine Learning

Spark is doing to ML on Hadoop what Hive did for SQL on Hadoop

Elegant Developer APIs

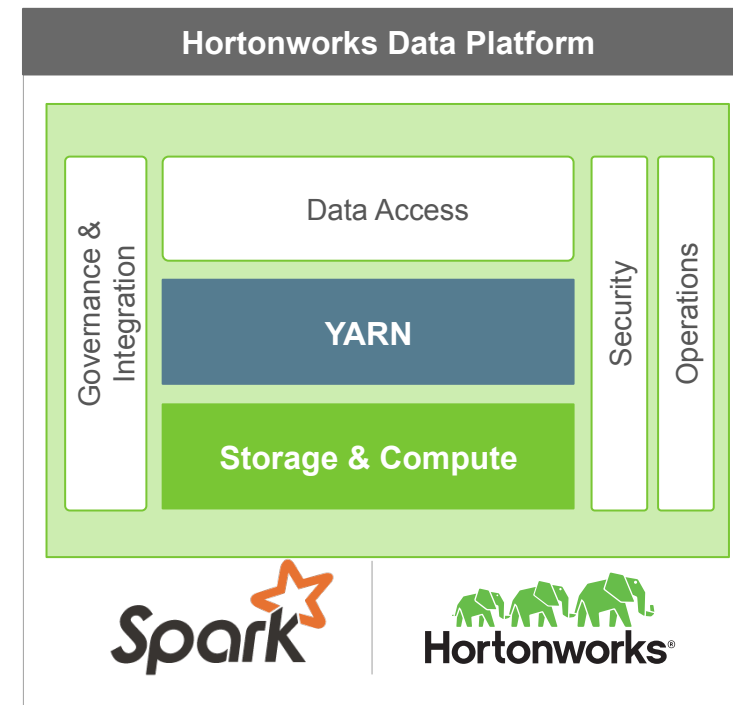
DataFrames, Machine Learning and SQL

Realize Value of Data Operating System

A key tool in the Hadoop toolbox

Community

Broad developer, customer and partner interest



Spark in Hadoop[®] with HDP 2.3

Resource Management

YARN for multi-tenant, diverse workloads with predictable SLAs

Tiered Memory Storage

HDFS in-memory tier—External BlockStore for RDD Cache

Deployable Anywhere

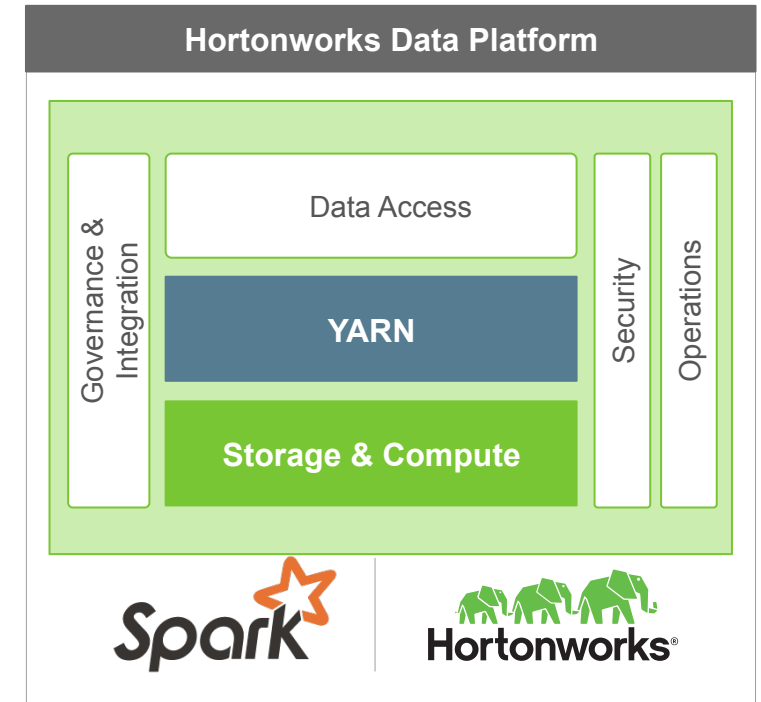
Linux, Windows and on-premises or cloud

Self-Service Spark in the Cloud

Easy launch of Data Science clusters via Cloudbreak and Ambari—for Azure, AWS, GCP, OpenStack and Docker

Operations

Deployment / management via Apache Ambari



Spark is Integrated into HDP's Centralized Architecture

HDP 2.3 Ships with Apache Spark 1.3.1

Production-ready

Centralized Resource Management

Run other workloads along with Spark
YARN provides capacity guarantees via
Capacity Scheduler

Consistent Operations


Deployable anywhere
Ambari deploys and manages

Comprehensive Security

Improved Authentication—only way to run
in a kerberized environment

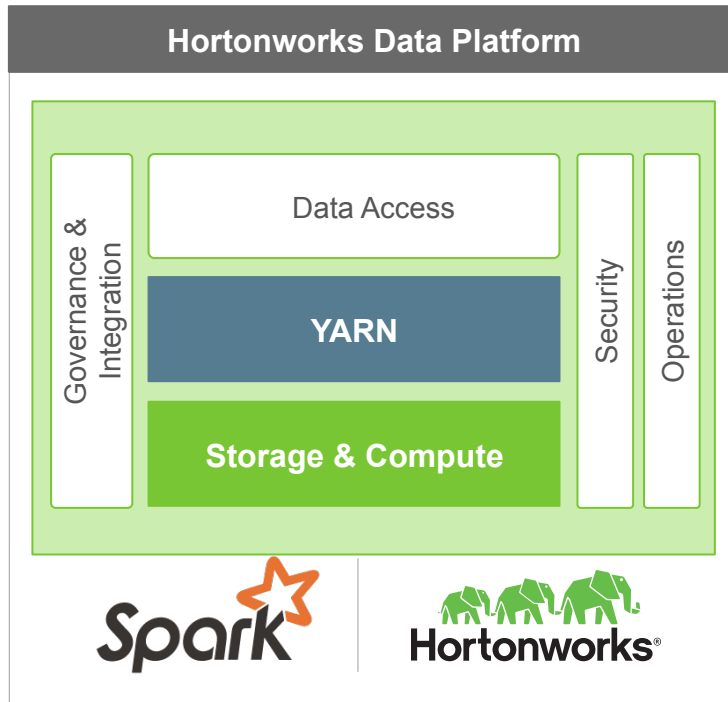
With Speed and at Scale

Vertical integration of Spark with YARN,
HDFS and ORC



* Tech Preview

Hortonworks Focus for Spark



Easy of Use

Apache Zeppelin for interactive notebooks

Metadata and Governance

Apache Atlas for metadata & Apache Falcon support for Spark pipelines

Security

Apache Ranger managed authorization

Spark SQL and Hive for SQL

Interop with modern Metastore / HS2, optimized ORC support, advanced analytics—e.g., Geospatial

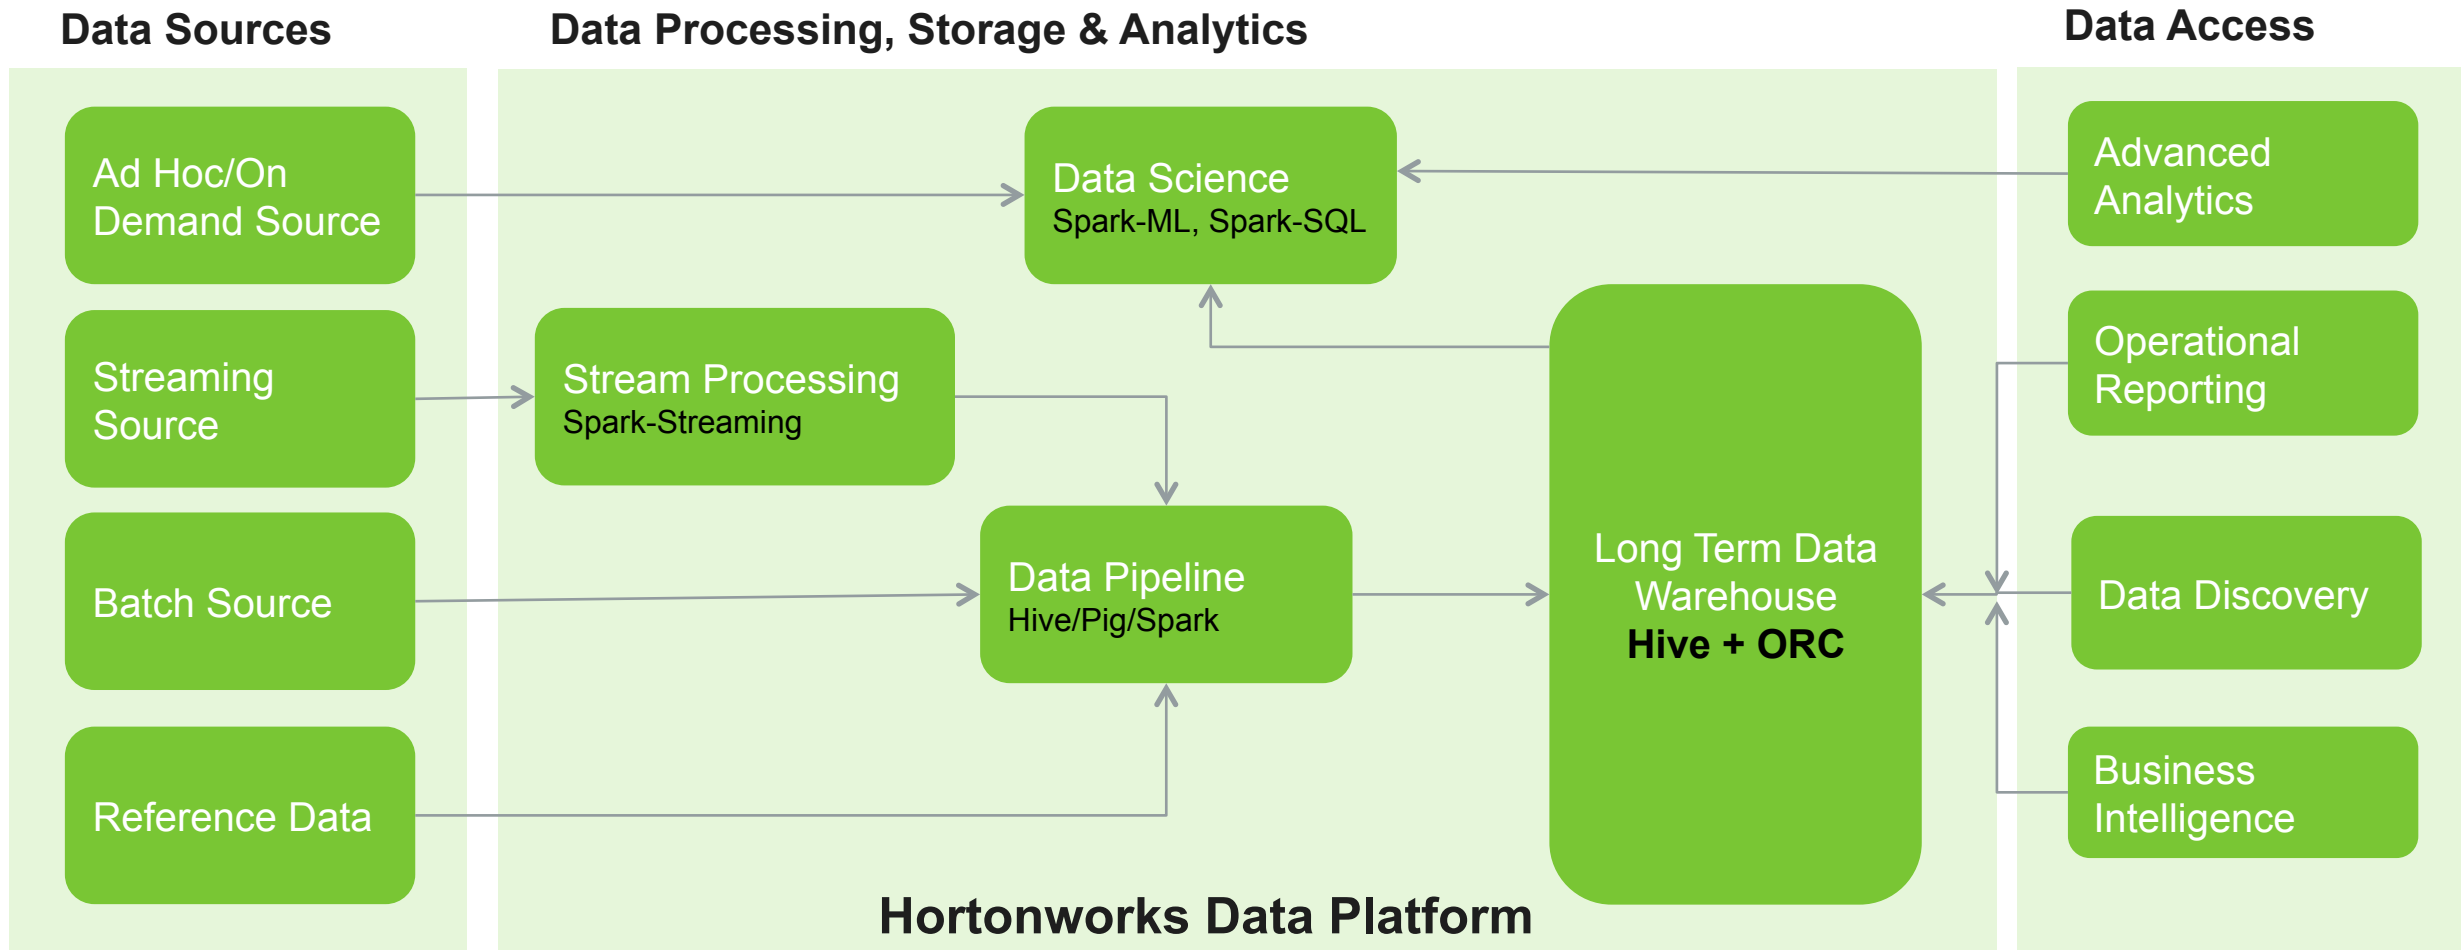
Spark and NoSQL

Deep integration with HBase via DataSources / Catalyst for Predicate / Aggregate Pushdown

Connect the Dots—Algorithms to Use-Cases

Higher-level ML Abstractions—e.g., OneVsRest
Validation, tuning, pipeline assembly—e.g., GeoSpatial

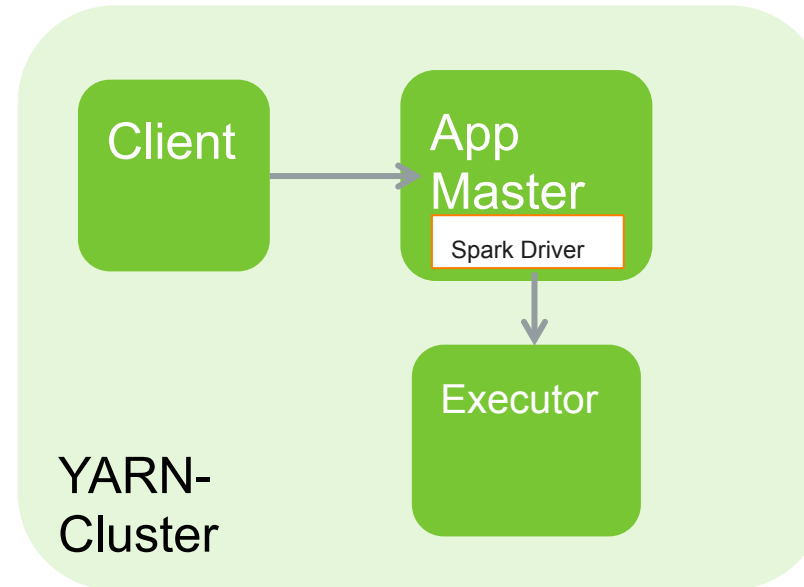
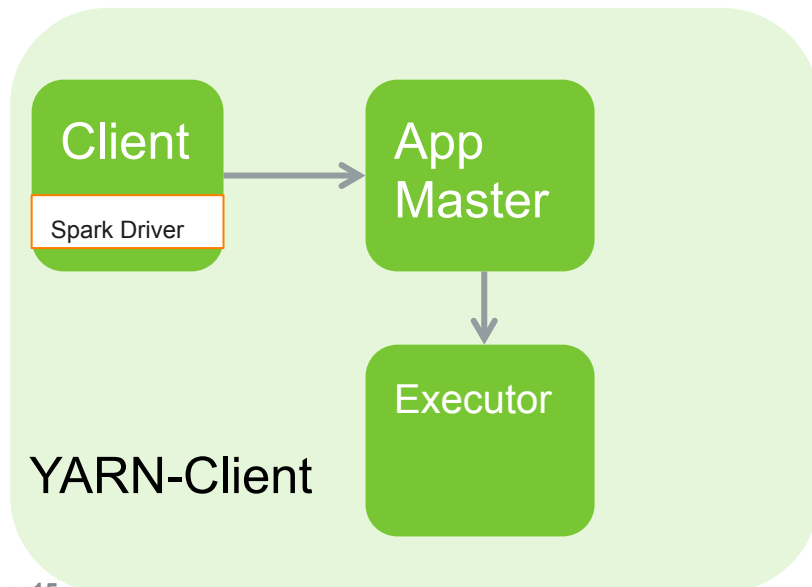
Reference Deployment Architecture



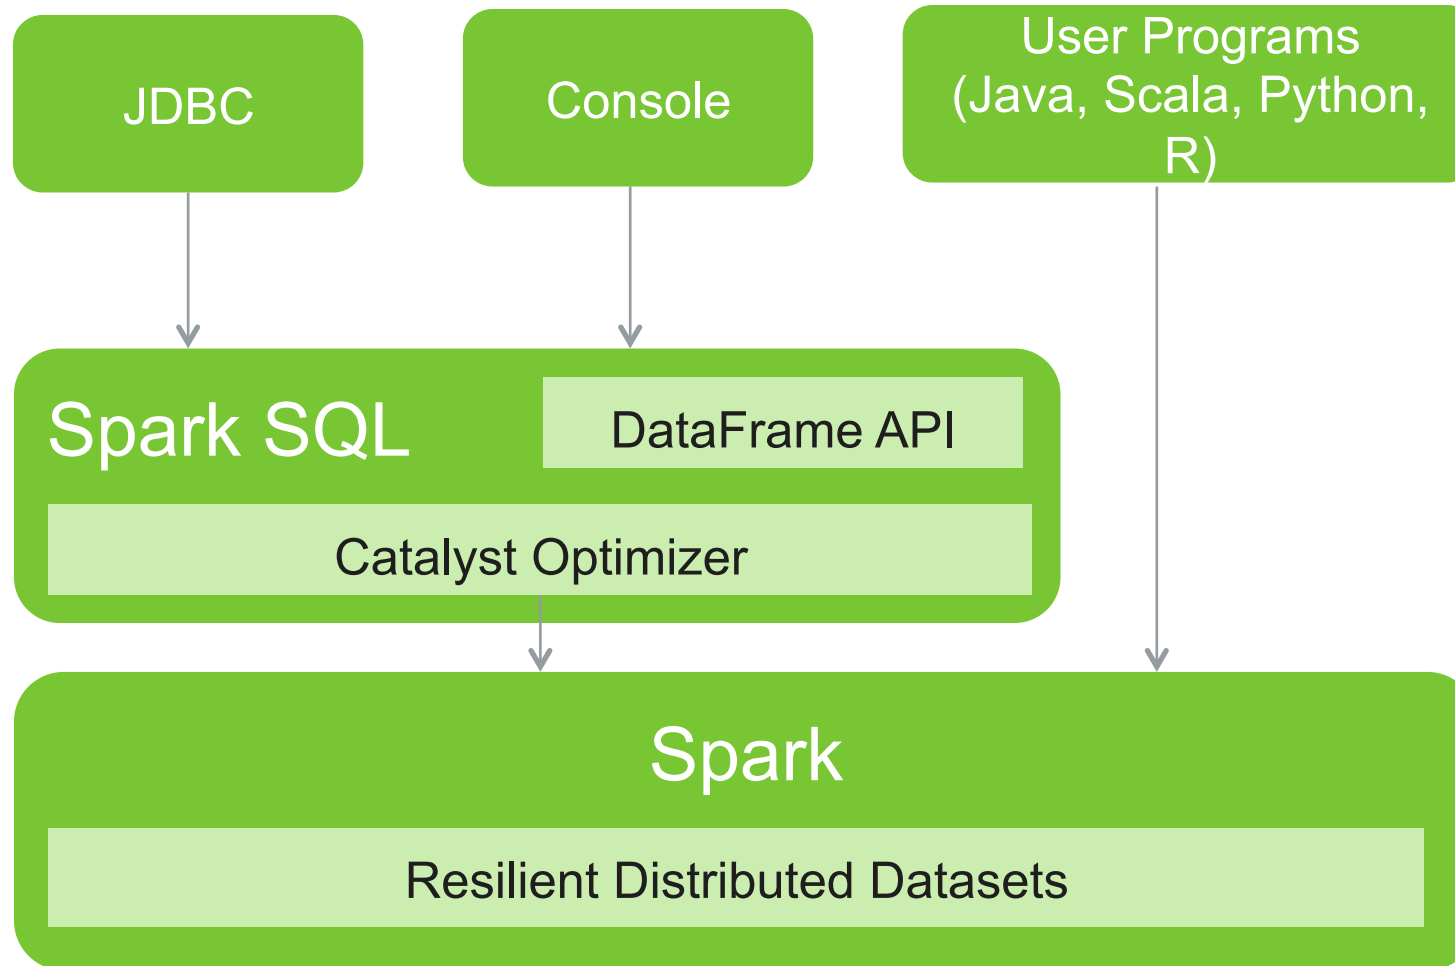
Spark Deployment Modes

- Spark Standalone Cluster
 - For developing Spark apps against a local Spark (similar to develop/deploying in IDE)
- Spark on YARN
 - Spark driver (SparkContext) in YARN AM(yarn-cluster)
 - Spark driver (SparkContext) in local (yarn-client)
 - Spark Shell runs in yarn-client only

Mode setup with
Ambari



Interfaces to Spark SQL



Current State of Security in Spark

HDP 2.3

Only Spark on YARN supports Kerberos today

Leverage Kerberos for authentication

Spark reads data from HDFS and ORC

HDFS file permissions (and Ranger integration) applicable to Spark jobs

Spark submits job to YARN queue

YARN queue ACL (and Ranger integration) applicable to Spark jobs

Wire Encryption

Spark has some coverage, not all channels are covered

LDAP Authentication

No authentication in Spark UI OOB, supports filter for hooking in LDAP

```
15/07/03 19:05:19 INFO AbstractConnector: Started SocketConnector@0.0.0.0:49270
15/07/03 19:05:19 INFO Utils: Successfully started service 'HTTP class server' on port 49270.
Welcome to
```

```

  _ _ _ _ _
 / _/ _ _ _/ _/
 \ V _ V _ ' / _/ ' \
/_/_/ . _ ^ ,/_/_/_/_/
  /_/_/

```

version 1.3.1

Using Scala version 2.10.4 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_45)

Type in expressions to have them

Type :help for more information.

15/07/03 19:05:21 INFO SparkCont

15/07/03 19:05:21 INFO SecurityM

15/07/03 19:05:21 INFO SecurityM

15/07/03 19:05:21 INFO SecurityM

```
fy permissions: Set(mLong)
```

Welcome to

```

      _ _ _ _ _
      /  _  \   _ _ _ _ _  /  _  \
     _\  V  _\  V  _\  /  _  \  '  /
    / _  / . _ ^ _ , _  _  / _  ^ \
      /  \

```

version 1.3.1

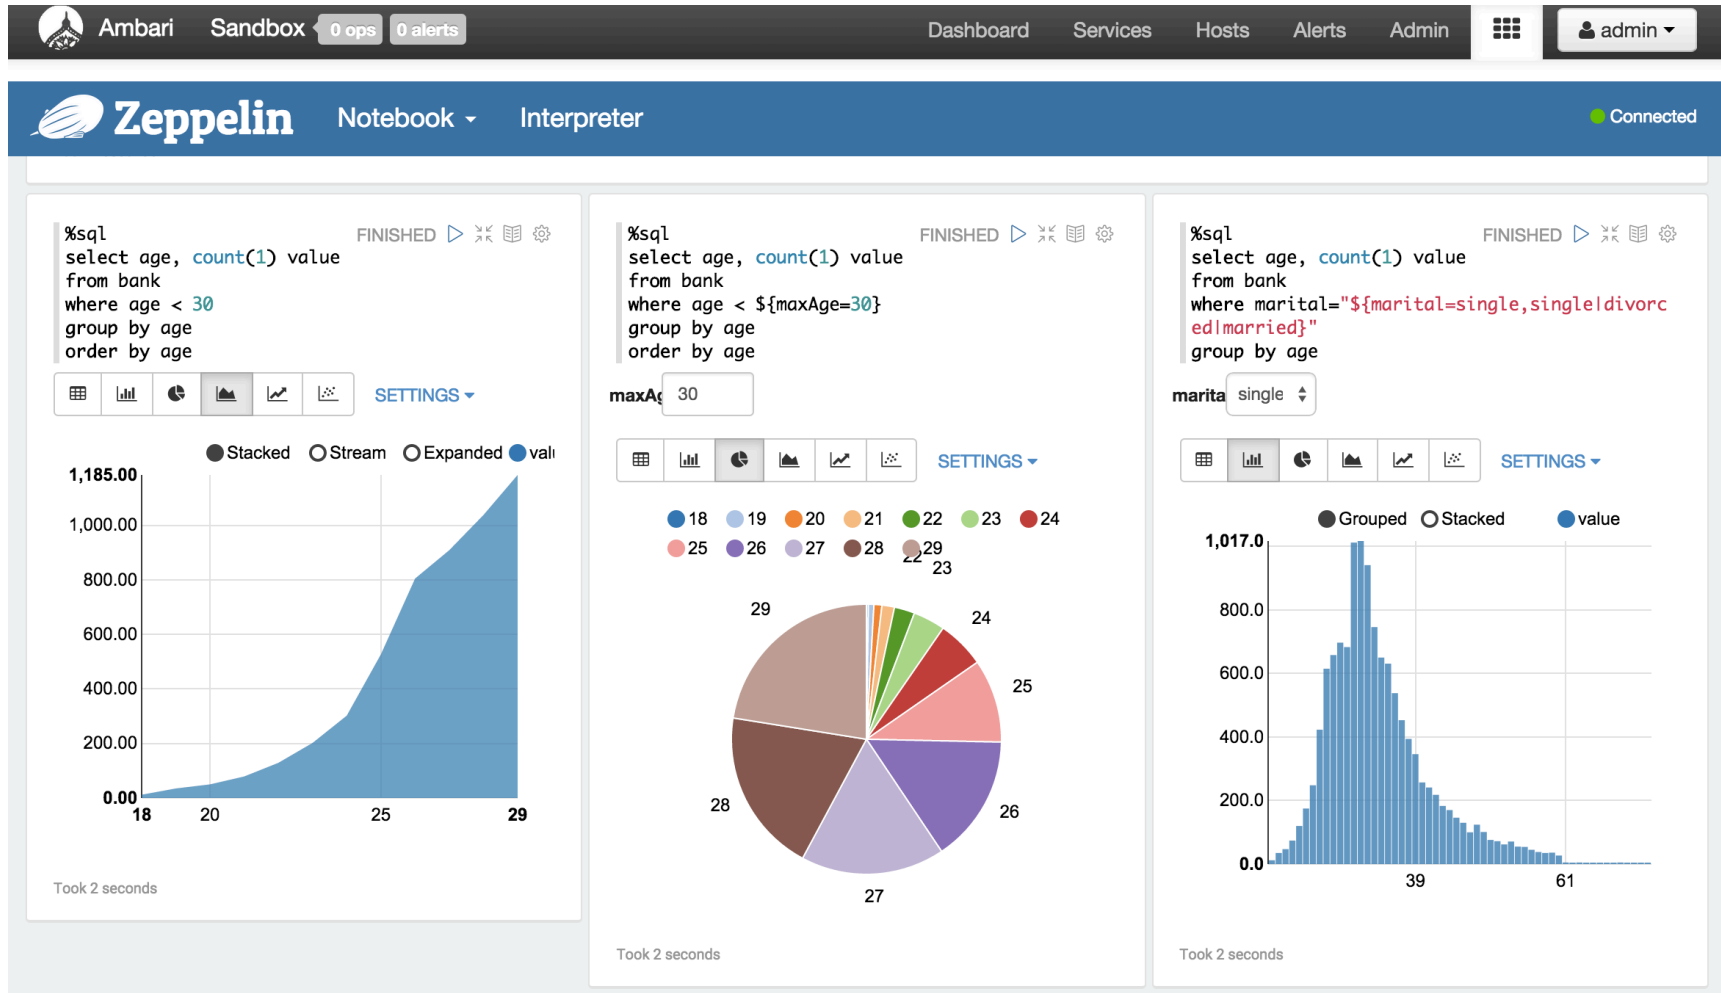
Using Python version 2.7.6 (default, Sep 9 2014 15:04:36)

SparkContext available as sc, HiveContext available as sqlContext.



Zeppelin

Introducing Apache Zeppelin



Apache Zeppelin

Features	Use Cases
<p>A web-based notebook for interactive analytics —Ad-hoc experimentation with Spark, Hive, Shell, Flink, Tajo, Ignite, Lens, etc</p> <p>Deeply integrated with Spark and Hadoop —Can be managed via Ambari Stacks</p> <p>Supports multiple language backends —Pluggable “Interpreters”</p> <p>Incubating at Apache —100% open source and open community</p>	<p>Data exploration and discovery</p> <p>Visualization—tables, graphs and charts</p> <p>Interactive snippet-at-a-time experience</p> <p>Collaboration and publishing</p> <p>“Modern Data Science Studio”</p>

Apache Zeppelin

```
HW11718:bin mlong$ ./zeppelin-daemon.sh start
```

```
Zeppelin start
```

[OK]

```
HW11718:bin mlong$
```



Zeppelin

Notebook ▾

Interpreter

salary demo



```
val salaryData = sc.textFile("data/salarydata.txt")
```

```
salaryData: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[58] at textFile at <console>:24
```

Took 0 seconds

```
val genderSalaryData = salaryData.map(line => line.split(',')).map(line => (line(0), line(2).toInt))
genderSalaryData.cache()
genderSalaryData.collect()
```

```
genderSalaryData: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[60] at map at <console>:26
```

```
res25: genderSalaryData.type = MapPartitionsRDD[60] at map at <console>:26
```

```
res26: Array[(String, Int)] = Array((M,39000), (F,41000), (M,99000), (M,58000), (M,43000), (M,11000), (M,0), (F,96000), (M,37000), (F,53000), (F,27000), (F,0), (M,54000), (F,0), (F,45000), (M,57000), (M,16000), (M,0), (M,0), (M,75000), (F,0), (F,42000), (F,48000), (F,16000), (F,85000), (F,72000), (M,18000), (M,8100), (F,69000), (F,57000), (M,76000), (M,12000), (M,58000), (F,96000), (F,37000), (F,0), (F,20000), (M,0), (F,0), (M,99000), (F,68000), (F,0)...
```

Took 0 seconds



PySpark / Spark SQL



Notebook ▾

Interpreter

```
%pyspark
py_val=sc.parallelize(["M",14,0,95102])
print py_val.collect()
```

['M', 14, 0, 95102]

Took 0 seconds



Zeppelin

Notebook ▾

Interpreter

```
%sql
select s.gender, s.salary from
(select d.gender, d.salary, row_number() over (partition by d.gender order by d.salary desc) rn
 from (select distinct gender, salary from salaries) d) s
where s.rn <= 10 order by gender, salary desc
```





Ambari

Sandbox

0 ops

0 alerts

Da



Zeppelin


Notebook ▾

Interpreter

```
%sql
select age, count(1) value
from bank
where age < 30
group by age
order by age
```

FINISHED ▶ ✕ 📖 ⚙️

SETTINGS ▾



Took 2 seconds


```
%sql
select age, count(1) value
from bank
where age < ${maxAge=30}
group by age
order by age
```

FINISHED

maxAge 30

ST

18 19 20 21 22
 25 26 27 28 29 23



Took 2 seconds

Web-based Notebook for interactive analytics

Features

Ad-hoc experimentation

Spark, Hive, Shell, Flink, Tajo, Ignite, Lens, etc

Deeply integrated with Spark + Hadoop

Can be managed via Ambari Stacks

Supports multiple language backends

Pluggable “Interpreters”

Incubating at Apache

100% open source and open community

Use Case

Data exploration and discovery

Visualization

tables, graphs and charts

Interactive snippet-at-a-time experience

Collaboration and publishing

“Modern Data Science Studio”

Demo – Parte 1

Spark Streaming

What is Spark Streaming?

- Extends Spark for doing large scale stream processing
 - Scales to 100s of nodes and achieves second scale latencies
- Efficient and fault-tolerant stateful stream processing
 - Simple batch-like API for implementing complex algorithms

Motivation


- Many important applications must process large streams of live data and provide results in near-real-time
 - Social network trends
 - Website statistics
 - Ad impressions
- Distributed stream processing framework is required to
 - Scale to large clusters (100s of machines)
 - Achieve low latency (few seconds)

Integration with Batch Processing

- Many environments require processing same data in live streaming as well as batch post processing
- Existing framework cannot do both
 - Either do stream processing of 100s of MB/s with low latency
 - Or do batch processing of TBs / PBs of data with high latency
- Extremely painful to maintain two different stacks
 - Different programming models
 - Double the implementation effort
 - Double the number of bugs

Stateful Stream Processing

- Traditional streaming systems have a **record-at-a-time** processing model
 - Each node has mutable state
 - For each record, update state and send new records
- State is lost if node dies.
- Making stateful stream processing be fault-tolerant is challenging



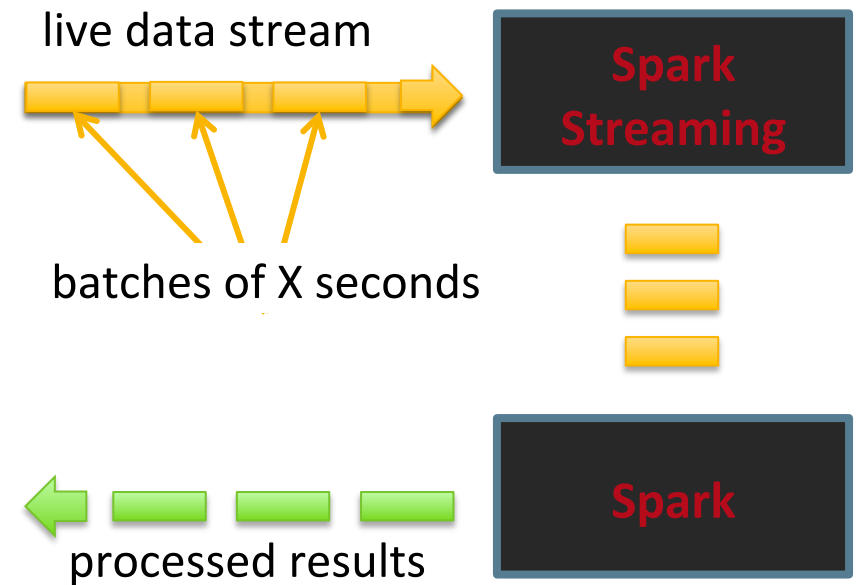
Comparison with Storm

- Storm
 - Replays record if not processed by a node
 - Processes each record *at least once*
 - May update mutable state twice
 - Mutable state can be lost due to failure
- Spark vs Storm: Use case dictates which one to use. Use Storm for at least once, and low latency SLAs. Can use Spark for SLAs > 500ms. Officially HWX doesn't support Spark Streaming and supports Storm.

Discretized Stream Processing

Run a streaming computation as a **series of very small, deterministic batch jobs**

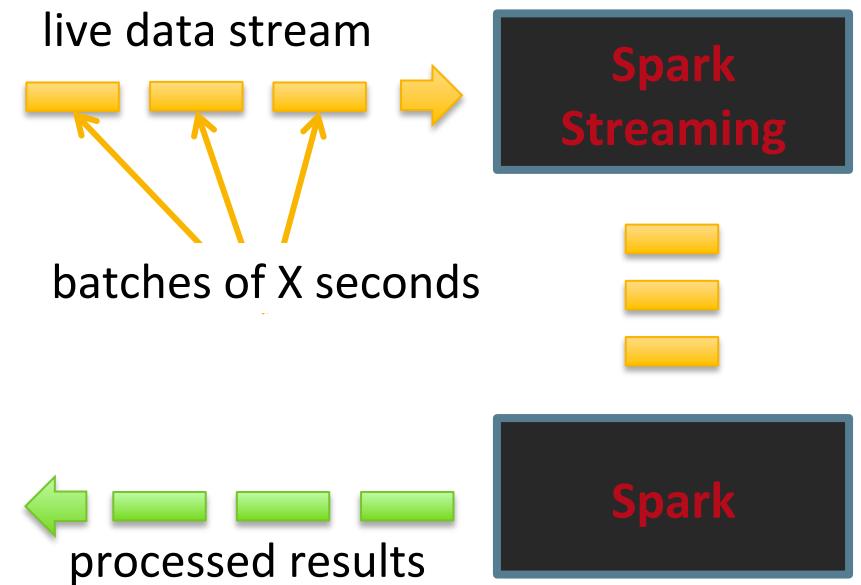
- Chop up the live stream into batches of X seconds
- Spark treats each batch of data as RDDs and processes them using RDD operations
- Finally, the processed results of the RDD operations are returned in batches



Discretized Stream Processing

Run a streaming computation as a **series of very small, deterministic batch jobs**

- Batch sizes as low as ½ second, latency of about 1 second
- Potential for combining batch processing and streaming processing in the same system



Example – Get hashtags from Twitter

```
val tweets = ssc.twitterStream()
```

DStream: a sequence of RDDs representing a stream of data

Twitter Streaming API

batch @ t

batch @ t+1

batch @ t+2



tweets DStream



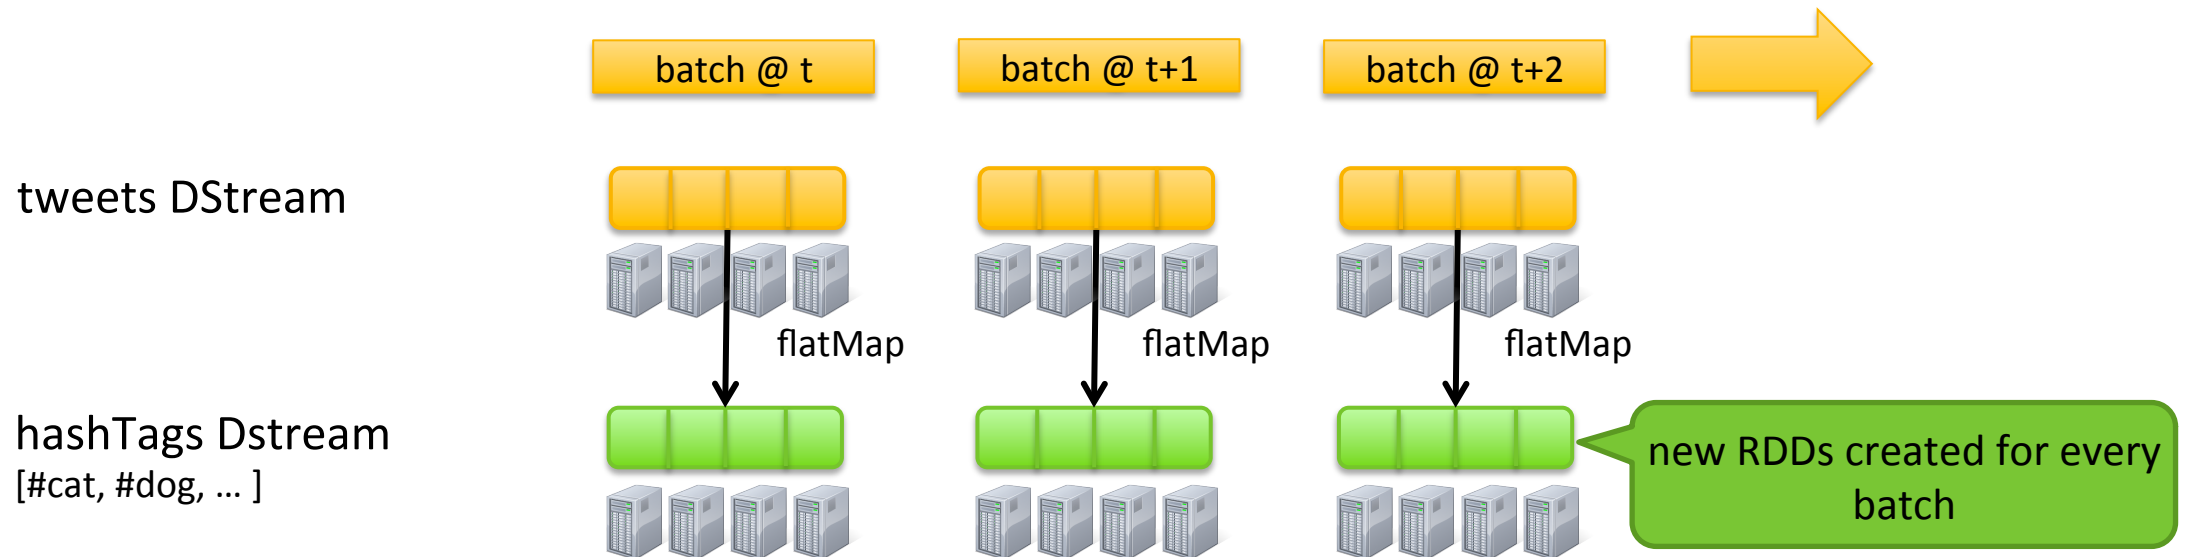
stored in memory as an RDD
(immutable, distributed)

Example – Get hashtags from Twitter

```
val tweets = ssc.twitterStream()  
val hashTags = tweets.flatMap (status => getTags(status))
```

new DStream

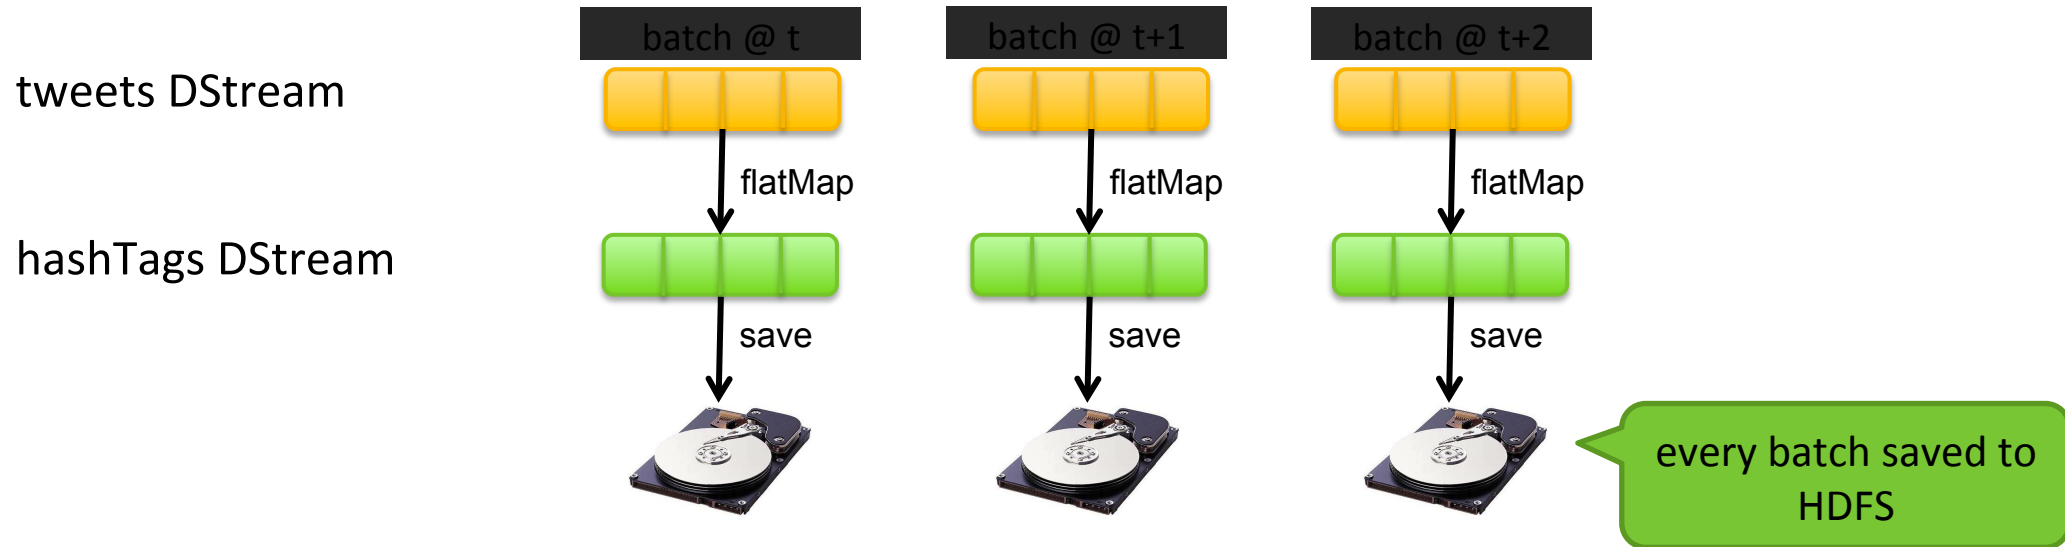
transformation: modify data in one DStream to create another DStream



Example – Get hashtags from Twitter

```
val tweets = ssc.twitterStream()  
val hashTags = tweets.flatMap (status => getTags(status))  
hashTags.saveAsHadoopFiles("hdfs://...")
```

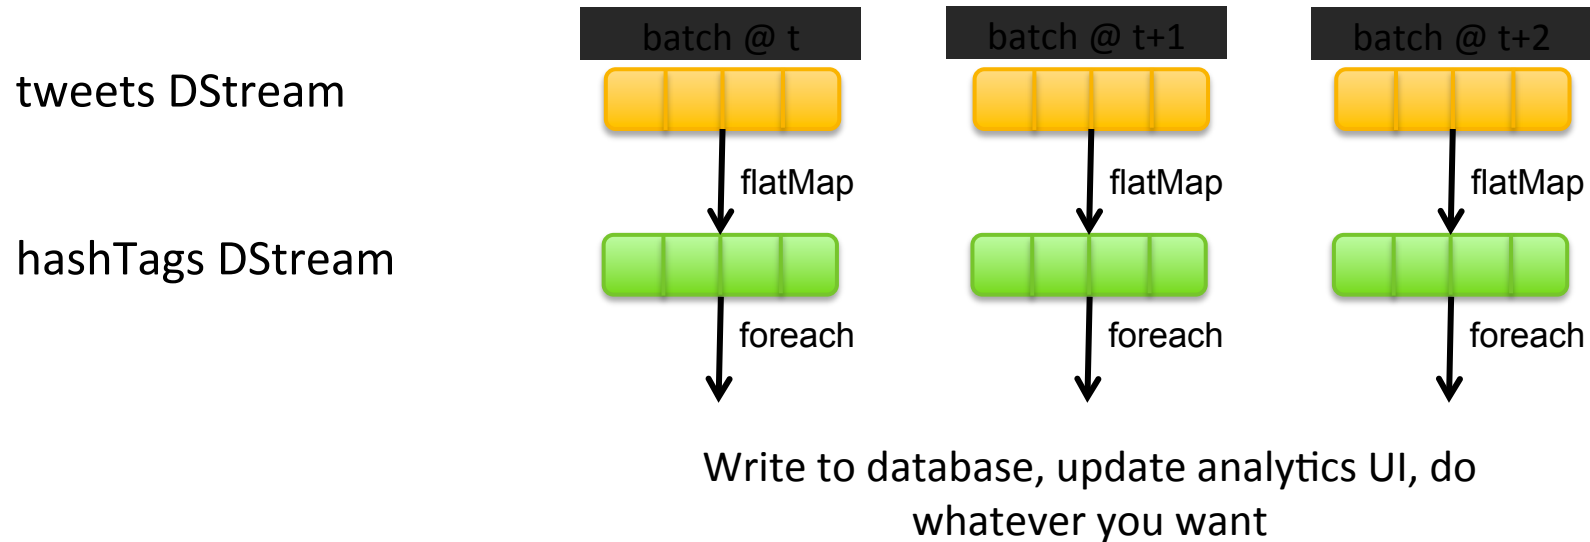
output operation: to push data to external storage



Example – Get hashtags from Twitter

```
val tweets = ssc.twitterStream()  
val hashTags = tweets.flatMap(status => getTags(status))  
hashTags.foreach(hashTagRDD => { ... })
```

foreach: do whatever you want with the processed data



Java Example

Scala

```
val tweets = ssc.twitterStream()  
val hashTags = tweets.flatMap(status => getTags(status))  
hashTags.saveAsHadoopFiles("hdfs://...")
```

Java

```
JavaDStream<Status> tweets = ssc.twitterStream()  
JavaDStream<String> hashTags = tweets.flatMap(new Function<...> { })  
hashTags.saveAsHadoopFiles("hdfs://...")
```

Function object


Window-based Transformations

```
val tweets = ssc.twitterStream()  
val hashTags = tweets.flatMap (status => getTags(status))  
val tagCounts = hashTags.window(Minutes(1), Seconds(5)).countByValue()
```

sliding window
operation

window length

sliding interval



Arbitrary Stateful Computations

Specify function to generate new state based on previous state and new data

- Example: Maintain per-user mood as state, and update it with their tweets

```
updateMood(newTweets, lastMood) => newMood  
moods = tweets.updateStateByKey(updateMood _)
```


Arbitrary Combinations of Batch and Streaming Computations

Inter-mix RDD and DStream operations!

- Example: Join incoming tweets with a spam HDFS file to filter out bad tweets


```
tweets.transform(tweetsRDD => {  
    tweetsRDD.join(spamHDFSFile).filter(...)  
})
```

DStream Input Sources

- Out of the box we provide
 - Kafka
 - HDFS
 - Flume
 - Akka Actors
 - Raw TCP sockets
- Very easy to write a *receiver* for your own data source

Fault-tolerance: Worker

- RDDs remember the operations that created them
- Batches of input data are replicated in memory for fault-tolerance
- Data lost due to worker failure, can be recomputed from replicated input data
- All transformed data is fault-tolerant, and exactly-once transformations



Fault-tolerance: Master

- Master saves the state of the DStreams to a checkpoint file
 - Checkpoint file saved to HDFS periodically
- If master fails, it can be restarted using the checkpoint file

Unifying Batch and Stream Processing Models

Spark program on Twitter log file using RDDs

```
val tweets = sc.hadoopFile("hdfs://...")  
val hashTags = tweets.flatMap(status => getTags(status))  
hashTags.saveAsHadoopFile("hdfs://...")
```

Spark Streaming program on Twitter stream using DStreams

```
val tweets = ssc.twitterStream()  
val hashTags = tweets.flatMap(status => getTags(status))  
hashTags.saveAsHadoopFiles("hdfs://...")
```

Demo – Parte 2

Apoio

Apoio

- ✓ Download Hortonworks Sandbox

<http://hortonworks.com/sandbox>

- ✓ Update Zeppelin Gallery

<https://github.com/hortonworks-gallery/zeppelin-notebooks>

Apoio

✓ Acessar Zeppelin

<http://localhost:9995>

✓ Acesso Shell

<http://localhost:4200>

✓ Outros Exemplos

<https://github.com/DhruvKumar/spark-twitter-sentiment>

<https://github.com/DhruvKumar/spark-workshop>

Apoio

✓ Twitter Apps

<https://apps.twitter.com/>

Perguntas?

Obrigado